

Kubeflow Control Plane

Weiqiang Zhuang

wzhuang@us.ibm.com

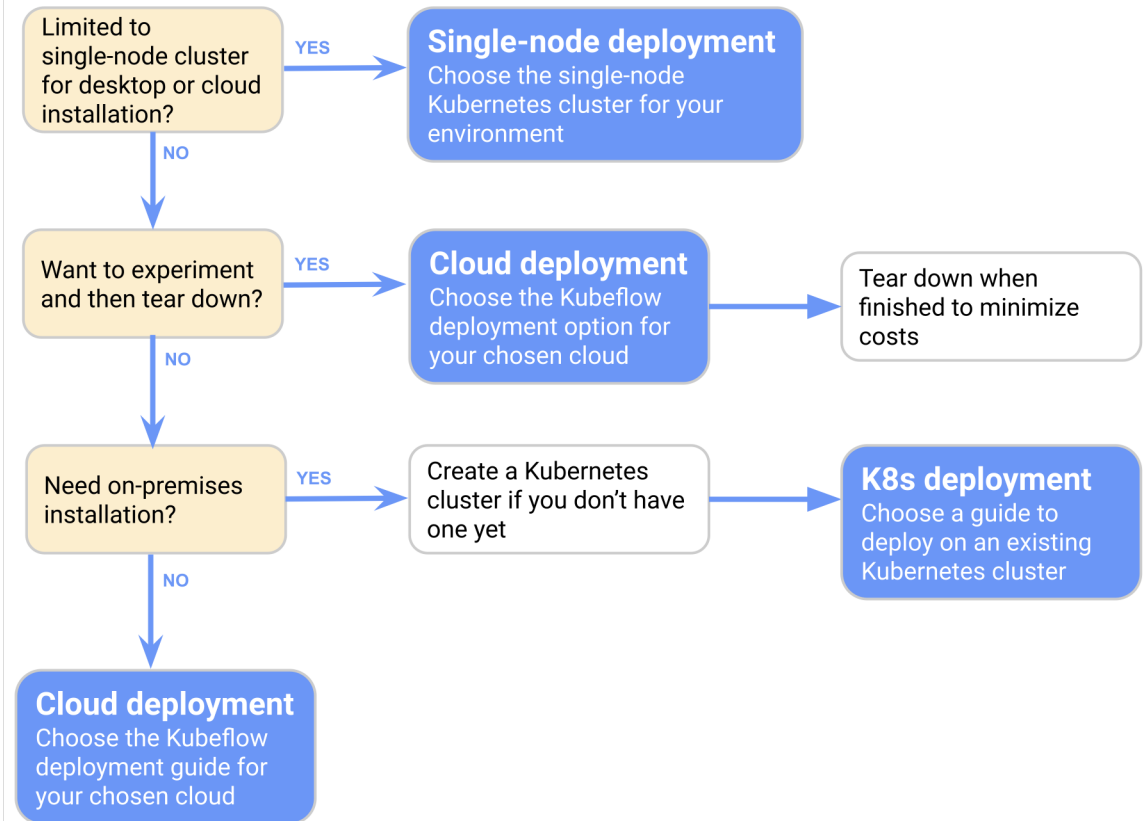
IBM CODAIT

Github id: adrian555



- Deploy, manage and monitor Kubeflow
 - Document
 - <https://www.kubeflow.org/docs/started/getting-started/>
 - On various environments
 - GCP/AWS/IKS/OpenShift
 - Other K8S
 - On-prem Linux/MacOS/Windows
 - minikube/miniKF
 - Deployed through command line or operator
 - Configuration for the collection of components/applications
 - Use one from [manifests](#) repo, or
 - Create your own
 - Two repos
 - kubectl <https://github.com/kubeflow/kfctl>
 - manifests <https://github.com/kubeflow/manifests>

Choosing a Kubeflow deployment option



<https://www.kubeflow.org/docs/images/kubeflow-getting-started-diagram.svg>



- *kfctl* is the control plane for deploying and managing Kubeflow.
 - Run *kfctl* as a CLI with **KfDef** configurations for different Kubernetes flavors
 - *kubeflow/kfctl* also incubates an [operator](#) to deploy and monitor Kubeflow
- **KfDef** configurations
 - are manifests specifying a set of applications to be deployed by their kustomization and resources by [kustomize](#)
 - resources of each application are organized in the layout for kustomize to process
 - *kubeflow/manifests* is the repo for the collection of KfDef configurations
- **kustomize**
 - customizes raw, template-free YAML files. It patches Kubernetes resources files with a kustomization file and various overlays.
 - kustomization is also a Kubernetes resource (kind: Kustomization). It contains the `generators` and `transformers` to be applied on the `resources`.



- KfDef

- Configuration through yaml

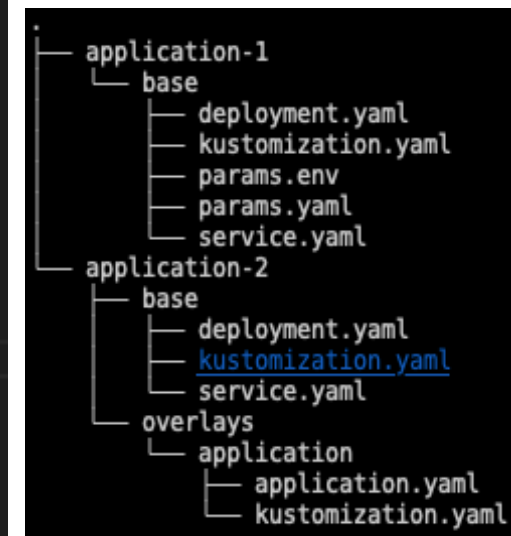
- Code

https://github.com/kubeflow/kfctl/blob/master/pkg/apis/apps/kfdef/v1/application_types.go

- *applications* are in kustomize form
- Also support plugins for certain platforms (ie. Aws, Gcp)
- Manifest repo can be either remote archive or local directory
 - The directory structure for manifests follows kustomize requirement
 - Eg. [Argo](#)

```
apiVersion: kfdef.apps.kubeflow.org/v1
kind: KfDef
metadata:
  name: kfdef-example
  namespace: kubeflow
spec:
  applications:
    - kustomizeConfig:
        name: application-1
        parameters:
          - name: param1
            value: value1
        repoRef:
          name: manifests
          path: application-1
    - kustomizeConfig:
        name: application-2
        overlays:
          - application
        repoRef:
          name: manifests
          path: application-2
  repos:
    - name: manifests
      url: https://example.com/manifests/v1.0.0.tar.gz
      version: v1.0.0
```

Configuration in yaml

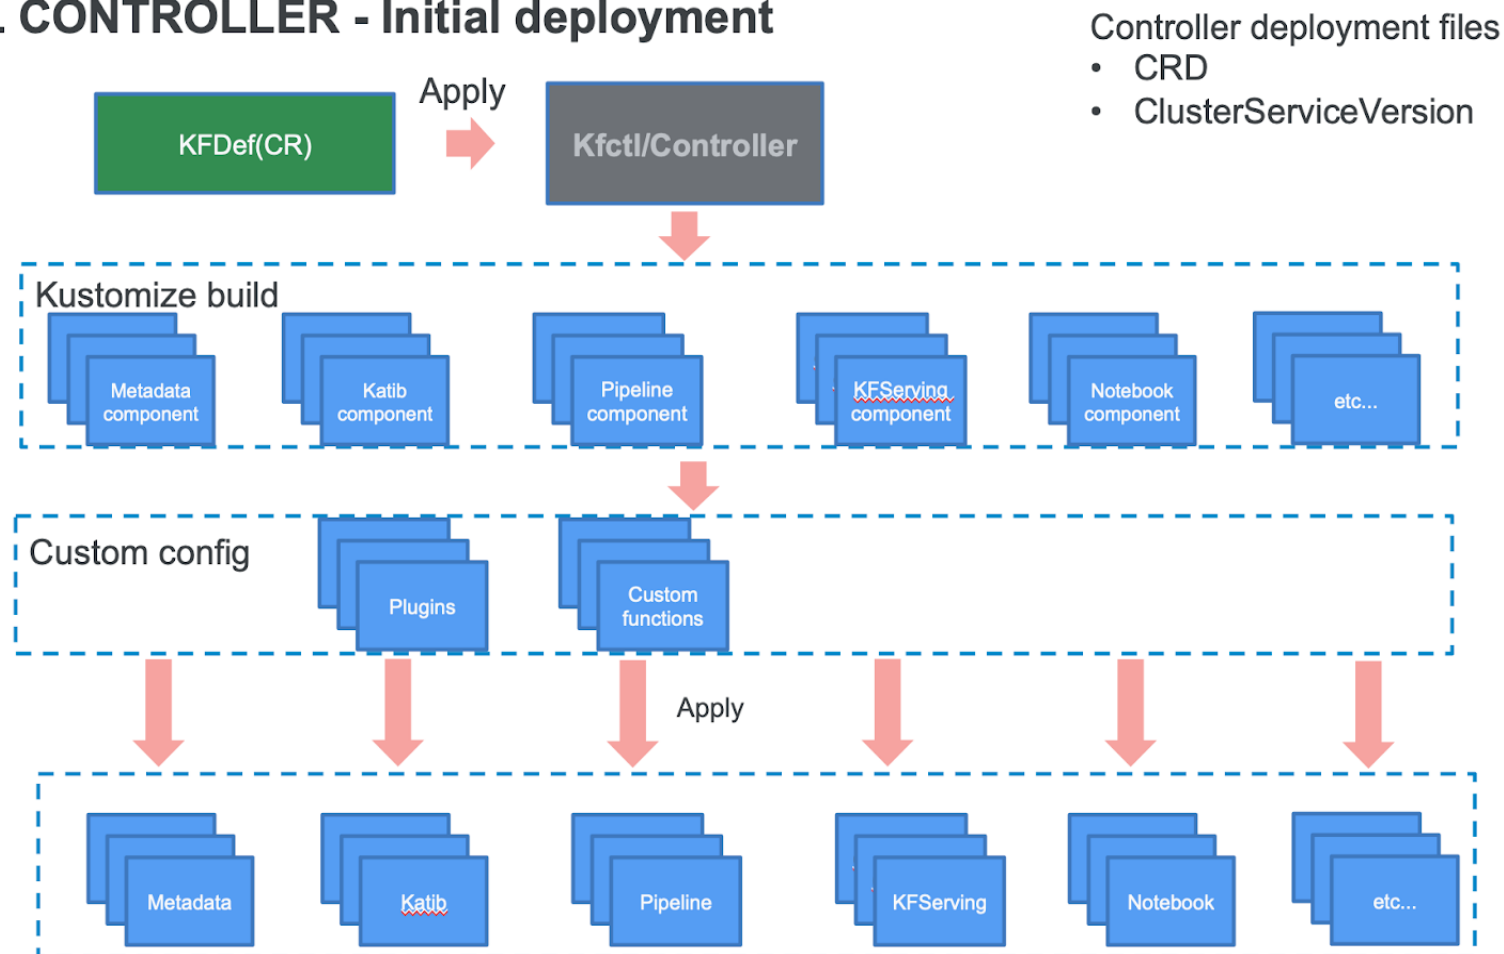


Directory structure



- kfctl

KFCTL CONTROLLER - Initial deployment



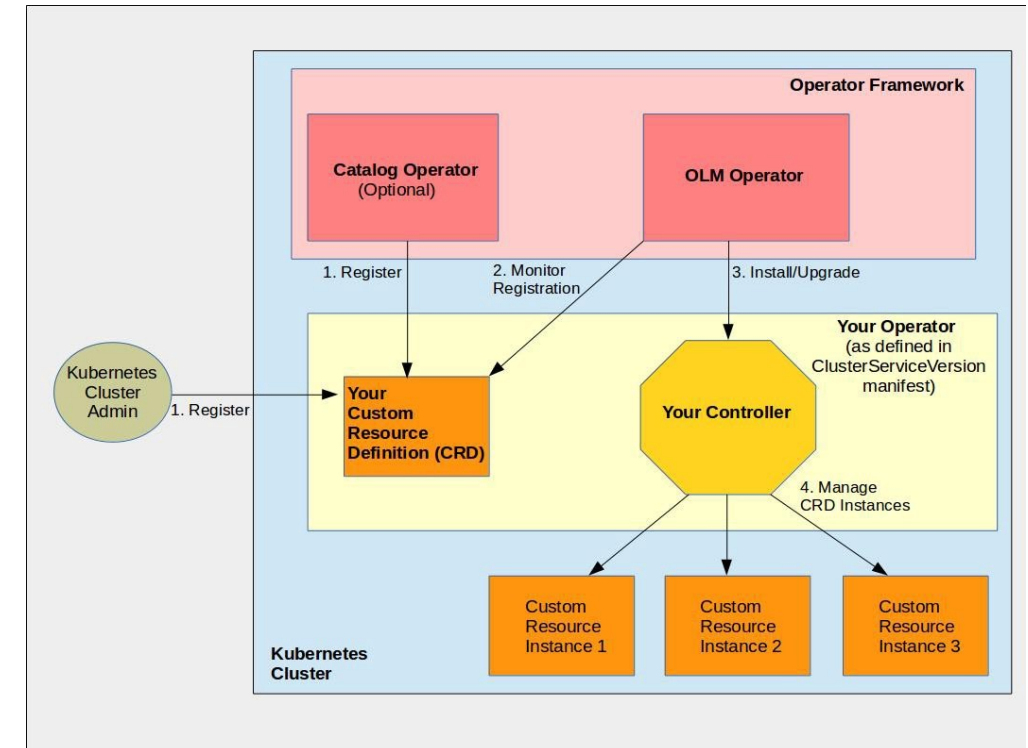
- `kfctl`
 - List all commands
 - `$> kfctl help`
 - Command line to install/uninstall Kubeflow
 - `$> kfctl build -f <config_uri>`
 - `$> kfctl apply -f <config_uri>`
 - `$> kfctl delete -f <config_uri>`
 - `<config_uri>` can be remote or local
 - Downloads the manifests for applications (if remote) from the *repo:uri* defined in the configuration file, and caches in the local disk
 - Loops through all applications' customization configuration and build/apply
 - Runs platform special handling if the configuration contains *plugins* section



- Kubeflow [manifests](#) repo
 - Maintains the manifests for Kubeflow's common applications
 - Argo, centraldashboard, admission-webhook, basic-auth, metadata, profiles and more
 - Other applications
 - Each application can be built with `kustomize` tool
 - `$> kustomize build`
 - `$> kubectl apply -k`



- Kubeflow operator
 - Operator helps deploy, monitor and manage the lifecycle of applications deployed on Kubernetes and OpenShift clusters
 - Built with [operator-sdk](#)
 - Learn more about operators - [link](#)
- Shares the same *apply* function with *kfctl* command
- *delete* function differs from *kfctl* command
- Configuration file is the custom resource (CR)
- Document
 - <https://github.com/kubeflow/kfctl/blob/master/operator.md>



https://miro.medium.com/max/2116/1*GYLAUB7KGCysjPgwek-pPA.jpeg



- Kubeflow operator
 - Code structure
 - [/deploy](#): Contains all the k8s resources for deploying the operator image and crd
 - [/build](#): Operator image build script
 - [/pkg/controller](#): main package for operator controller logic
 - [/cmd/manager](#): main.go file for the operator go program
 - Kubeflow operator watches the KfDef and other related resources
 - Two steps to install Kubeflow
 - Deploy the Kubeflow operator, then
 - Install the Kubeflow by creating the KfDef CR
 - Kubeflow operator continues to monitor and manage any KfDef CR created



- Kubeflow operator
 - Operator can be deployed by command line
 - `$> export OPERATOR_NAMESPACE=operators`
 - `$> kubectl create ns ${OPERATOR_NAMESPACE}`
 - `$> cd deploy/`
 - `$> kustomize edit set namespace ${OPERATOR_NAMESPACE}`
 - `$> kustomize build | kubectl apply -f -`
 - Operator is registered on operatorhub.io, can be installed through OLM console
 - OLM discovers the Kubeflow operator from its catalog source
 - Kubeflow is installed either by command lines or by subscription
 - creating a KfDef CR from command line
 - download the KfDef configuration file from `kubeflow/manifests`
 - add `metadata.name`
 - `$> kubectl apply -f <kfdef_configuration.yaml>`
 - creating a *subscription* to the operator from the OLM console



- More topics
 - KfUpgrade
 - other kfctl sub-commands
 - `kpt fn` commands
 - customize v3 support in the coming release 1.1
 - code walkthrough



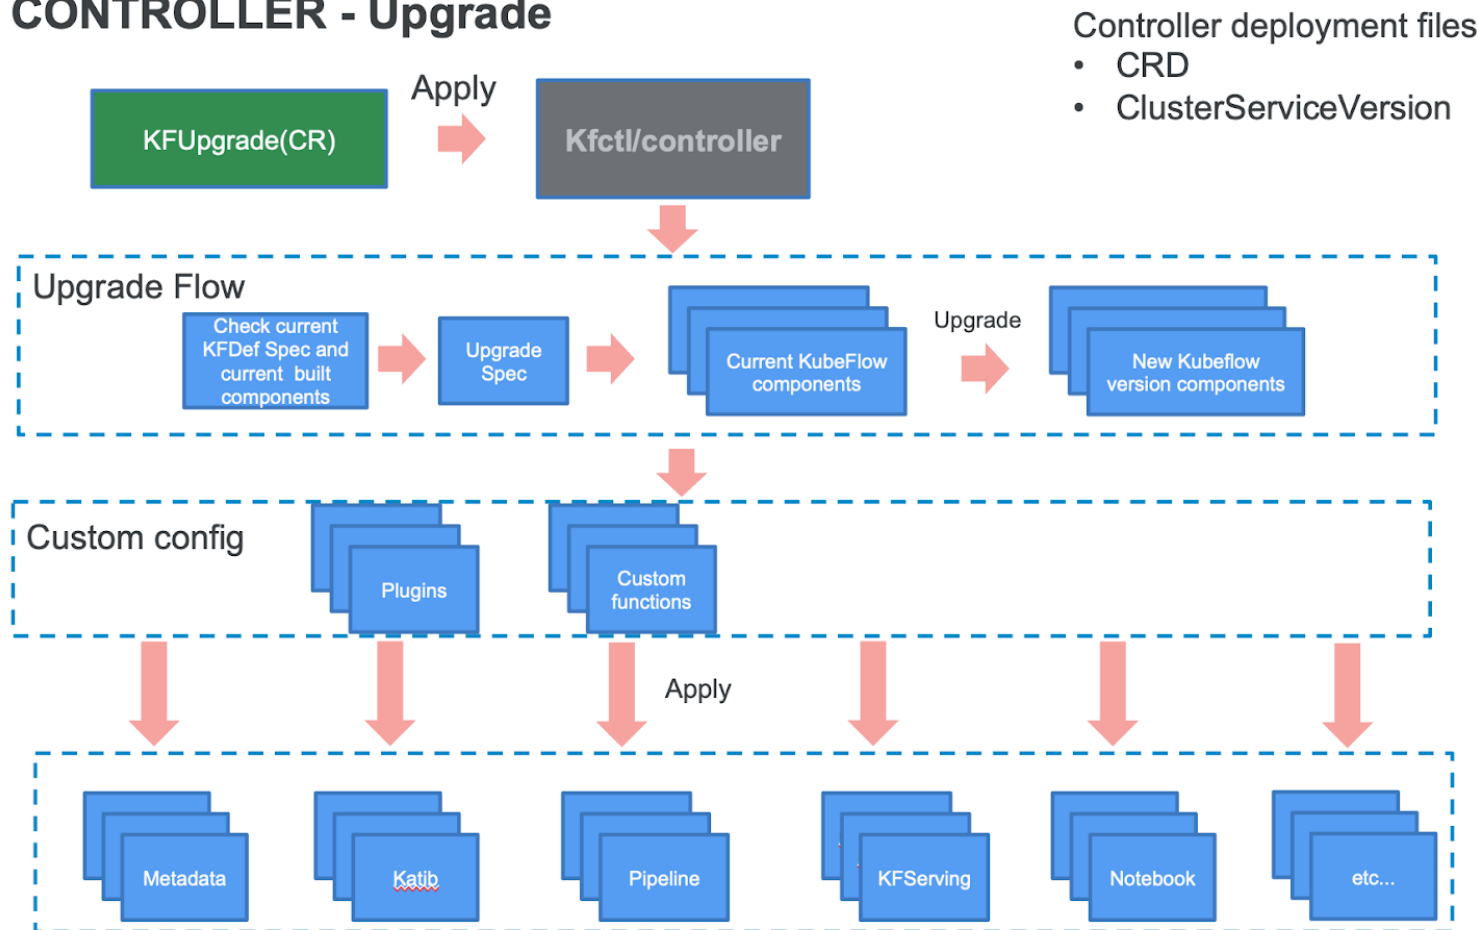
- KfUpgrade
 - in alpha

```
apiVersion: kfupgrade.apps.kubeflow.org/v1alpha1
kind: KfUpgrade
metadata:
  name: kf-upgrade-v0.7.1
spec:
  currentKfDef:
    # Replace with the name of your Kubeflow app
    name: kubeflow-app
    version: v0.7.0
  newKfDef:
    # Replace with the name of your kubeflow app
    name: kubeflow-app
    version: v0.7.1
  # Replace this with the path to the KfDef that you are upgrading to
  baseConfigPath: https://example.com/manifests/v0.7.1.yaml
```



- kfctl

KFCTL CONTROLLER - Upgrade



- Directories

```
kfctl
├── build
│   └── bin
├── cmd
│   ├── kfctl
│   │   └── cmd
│   ├── manager
│   └── plugins
│       └── dockerfordesktop
├── config
├── deploy
│   ├── crds
│   ├── olm-catalog
│   └── kubeflow
├── hack
├── pkg
│   ├── apis
│   │   └── apps
│   ├── controller
│   │   └── kfdef
│   ├── kfapp
│   │   ├── aws
│   │   ├── coordinator
│   │   ├── dockerfordesktop
│   │   ├── existing_arrikto
│   │   ├── gcp
│   │   ├── kustomize
│   │   ├── minikube
│   ├── kfconfig
│   │   ├── awsplugin
│   │   ├── gcpplugin
│   │   ├── loaders
│   │   └── testdata
│   ├── kfulgrade
│   │   └── mirror
│   │       └── testdata
│   └── utils
├── py
│   ├── kubeflow
│   └── kfctl
├── testing
│   └── workflows
│       ├── components
│       ├── environments
│       └── lib
└── third_party
```

```
manifests
├── admission-webhook
├── application
├── argo
├── aws
├── cert-manager
├── common
├── default-install
├── dex-auth
├── docs
├── experimental
├── gatekeeper
├── gcp
├── hack
├── istio
├── istio-1-3-1
├── jupyter
├── katib
├── kfdef
├── kfserving
├── knative
├── kubebench
├── kubeflow-roles
├── metacontroller
├── metadata
├── modeldb
├── mpi-job
├── mxnet-job
├── pipeline
├── plugins
├── profiles
├── pytorch-job
├── seldon
├── spark
├── tektoncd
├── tensorboard
├── tests
└── tf-training
```

```
kfdef
├── OWNERS
├── README.md
├── generic
│   ├── OWNERS
│   ├── README.md
│   ├── auth_oidc
│   │   ├── authservice.tmpl
│   │   ├── dex.tmpl
│   │   ├── envoy-filter.yaml
│   │   └── gateway.yaml
│   └── istio
│       ├── crds.yaml
│       └── istio-noauth.yaml
├── kfctl_anthos.v1.0.0.yaml
├── kfctl_anthos.v1.0.1.yaml
├── kfctl_anthos.yaml
├── kfctl_aws.v1.0.0.yaml
├── kfctl_aws.v1.0.1.yaml
├── kfctl_aws.yaml
├── kfctl_aws_cognito.v1.0.0.yaml
├── kfctl_aws_cognito.v1.0.1.yaml
├── kfctl_aws_cognito.yaml
├── kfctl_gcp_asm_exp.yaml
├── kfctl_gcp_basic_auth.v1.0.0.yaml
├── kfctl_gcp_basic_auth.v1.0.1.yaml
├── kfctl_gcp_basic_auth.yaml
├── kfctl_gcp_iap.v1.0.0.yaml
├── kfctl_gcp_iap.v1.0.1.yaml
├── kfctl_gcp_iap.yaml
├── kfctl_ibm.v1.0.0.yaml
├── kfctl_ibm.v1.0.1.yaml
├── kfctl_ibm.yaml
├── kfctl_istio_dex.v1.0.0.yaml
├── kfctl_istio_dex.v1.0.1.yaml
├── kfctl_istio_dex.yaml
├── kfctl_k8s_istio.v1.0.0.yaml
├── kfctl_k8s_istio.v1.0.1.yaml
├── kfctl_k8s_istio.yaml
└── kfctl_upgrade_gcp_iap_1.0.0.yaml
```

```
kfserving
├── kfserving-crds
│   ├── base
│   │   ├── crd.yaml
│   │   └── kustomization.yaml
│   └── overlays
│       └── application
│           ├── application.yaml
│           └── kustomization.yaml
└── kfserving-install
    ├── base
    │   ├── cluster-role-binding.yaml
    │   ├── cluster-role.yaml
    │   ├── config-map.yaml
    │   ├── kustomization.yaml
    │   ├── params.env
    │   ├── params.yaml
    │   ├── secret.yaml
    │   ├── service.yaml
    │   └── statefulset.yaml
    └── overlays
        └── application
            ├── application.yaml
            └── kustomization.yaml
```

